

NATURAL LANGUAGE PROCESSING



WORDS AND SENTENCES

What is a Word?

Easy Question?



What are words?

- “This is my cat, I call her Kitty.”
 - This sentence has 8 words.
 - How do you treat punctuation (here: “,” and “.”)? If you count it as words the sentence will have 10 words.
- Now consider speech applications. You have a transcript of spoken language:
 - “I um do large- largely data uh processing work.”
 - How do you treat filled pauses like “um” or “uh” sounds? Are these words?
 - How do you treat fragments like “large-”? Is it counted as a separate word?
- For speech applications there may be good reasons to keep fragments or filled pauses or to filter them out.

What are words?

- “This is my cat, I call her Kitty. Do you have cats, too?”
 - Here we have the words “cat” and “cats”
 - Both have the same word stem “cat”, also called lemma.
 - The two words have the same lemma but different wordform
 - Examples
 - Nouns in singular and plural have the same lemma but typically different wordforms
 - Verbs depending on time or grammatical form: show, showing, showed

How Would You Identify a Sentence?



Practical Part

- Go back to your notebook NLP01 and identify all sentences in The Time Machine.
- Do not use sentence tokenizers from any packages, yet.
- Implement your ideas.
- Keep a record of your sentences.
- How many sentences did you identify?

Organization of Texts

- In NLP a piece of text is usually called a document
- It contains words organized in sentences.
 - A sentence is usually a series of words finished by a characteristic punctuation: fullstop (“.”), question mark (“?”), exclamation mark (“!”)
 - Difficult to treat the colon (“:”)
 - “Then she said: ‘I never come back.’ → Two sentences
 - “He wrote his list: milk, honey, eggs, bread.” → one sentence.
 - However, exceptions exist:
 - Not every fullstop marks the end of a sentence, e.g. “e.g.” or “H.G. Wells”)
 - Some sentences do not end with “.” but e.g. with certain kinds of whitespaces. Typical example: Headlines

Organization of Texts

- Splitting according to punctuation needs to consider exceptions:
 - Not every fullstop marks the end of a sentence, e.g. “e.g.” or “H.G. Wells”)
 - IP addresses 10.0.0.1
 - Prices \$39.99
 - Some sentences do not end with “.” but e.g. with certain kinds of whitespaces. Typical example: Headlines

Text Processing

- Before we can analyze any text we usually need the following preprocessing steps:
 - Sentence tokenization (splitting the text into sentences)
 - Word tokenization (splitting the text into words)
 - Word normalization (reducing the words to their corresponding lemmata)

- Before we can analyze any text we usually need the following preprocessing steps:
 - **Sentence tokenization** (splitting the text into sentences)
 - Word tokenization (splitting the text into words)
 - Word normalization (reducing the words to their corresponding lemmata)

Sentence Tokenization

- We have explored regular expressions before.
- Regular expression could be used to subdivide a document into sentences.
- However, you have to account for all the exceptions mentioned before...
- ... and those we haven't considered yet...
- Good sentence tokenization requires more sophisticated methods.

Unsupervised Sentence Tokenization Algorithm

- T. Kiss and J. Strunk, Unsupervised Multilingual Sentence Boundary Detection. Computational Linguistics 32: 485-525 (2006)
- Paper: <https://aclanthology.org/J06-4003.pdf>
- The algorithm is based on classification of occurrences of a fullstop as
 - Abbreviation (street “St.”, abbreviation “abbrev.”, for example “e.g.”, ...)
 - Ellipsis (...)
 - Initials (H.G. Wells)
 - Ordinal numbers (1. Chapter, I. Chapter)
- Special case: Abbreviation at the end of a sentence (“I live at Porter St.”)

Sentence Boundary Detection

Idea

- The algorithm classifies the occurrences of fullstops in several steps
 - First, classify into abbreviation, ellipsis, sentence end.
- In the second stage these finds are refined:
 - Abbreviations are checked whether they coincide with a sentence end, being reclassified as sentence end.
 - Ellipses are checked whether they coincide with a sentence end, being reclassified as sentence end.
 - Checking whether a sentence end is actually an initial, being reclassified as abbreviation.
 - Checking whether a sentence end is actually an ordinal number, being reclassified as abbreviation.

Sentence Boundary Detection

Likelihood Ratios

- The algorithm is based on likelihood ratios
 - How likely is a point to appear behind a certain word.
 - Normal words appear in any part of a sentence and usually not only at the sentence end.
 - Therefore, the likelihood ratio of word followed by fullstop / word not followed by fullstop is small.
 - Abbreviations are certain to be followed by a fullstop.

Sentence Boundary Detection

Heuristics

- In the second classification stage three heuristics are used
 - orthographic heuristic
 - collocation heuristic
 - frequent sentence starter heuristic

Sentence Boundary Detection

Heuristics

- Orthographic heuristic
 - After a sentence end, the following word is written capitalized.
 - Problems:
 - Some words are usually written capitalized independent of a sentence ending, e.g. names, “I”, ...
 - Some words are never written capitalized (e.g. mathematical equations)
 - Use likelihood ratios: Count whether the word in question is written at least once in lower case.
 - Example: “I don’t like abbrev. because they make trouble. Peter likes abbrev. They make his life easier.”
 - In this sample the first occurrence of abbrev. is in the middle of a sentence. The second one at the end of a sentence.
 - They is written capitalized in the second case indicating a sentence end.
 - This is valid since “they” occurs at least once uncapitalized, therefore is not always written capitalized.

Sentence Boundary Detection

Heuristics

- Collocation heuristic
 - Abbreviations often show collocations left and right of the fullstop
 - E.g. “e.g.”, “i.e.”, “et. al.”
 - Sentence boundary typically do not show collocation.
 - If there is a high likelihood ratio for the words left and right of the fullstop it is an abbreviation rather than a sentence boundary.

Sentence Boundary Detection

Heuristics

- Frequent sentence starter heuristic
 - Determine words which frequently occur at the beginning of a sentence.
 - Use part of the text where the sentence boundary is sure.
- For benchmarks and further details, please read the original paper.

- Before we can analyze any text we usually need the following preprocessing steps:
 - Sentence tokenization (splitting the text into sentences)
 - **Word tokenization** (splitting the text into words)
 - Word normalization (reducing the words to their corresponding lemmata)

Word Tokenization

- Basically, word tokenization means splitting text into single words.
- Can be done pattern based using regular expressions.
- Typical pattern for a word `[A-Za-z]+`
- Typical patterns indicating the end of a word:
 - Punctuation
 - End of string
 - Whitespaces

Word tokenization

- Typical problems in word tokenization:
 - A period is not necessarily the end of a sentence
 - IP addresses 127.0.0.1
 - Dates (in Europe) 10.10.2023
 - Decimal numbers (in English) \$ 9.99
 - Ellipsis
 - Clitic contractions
 - don't, doesn't = do not
 - I'll = I will
 - ...

- The Python-based NLP toolkit NLTK offers several word tokenizers.
- `nltk.regexp_tokenize()`

```
>>> text = 'That U.S.A. poster-print costs $12.40...'
>>> pattern = r'''(?x)      # set flag to allow verbose regexps
...     (?:[A-Z]\.)+        # abbreviations, e.g. U.S.A.
...     | \w+(?:-\w+)*       # words with optional internal hyphens
...     | \$?\d+(?:\.\d+)?%? # currency and percentages, e.g. $12.40, 82%
...     | \.\.\.            # ellipsis
...     | [][.,;"'()?():-_\` ] # these are separate tokens; includes ], [
...
>>> nltk.regexp_tokenize(text, pattern)
['That', 'U.S.A.', 'poster-print', 'costs', '$12.40', '...']
```

Code from Bird et. al., NLTK Book, Chapter 3

- Before we can analyze any text we usually need the following preprocessing steps:
 - Sentence tokenization (splitting the text into sentences)
 - Word tokenization (splitting the text into words)
 - **Word normalization** (reducing the words to their corresponding lemmata)

Word Normalization

- **Normalization** → Putting into standard format representing multiple forms
 - **Lemmatization** → Putting each word to its root, for example
 - am, are, is → be (irregular verbs)
 - Word, words → word (plural forms)
 - **Case folding** → putting everything in a single way of use of capital letters, for example
 - The Time Machine → the time machine (all to lower case)

Lemmatization Algorithms

- Most lemmatization is based on decomposing words into parts (**morphemes**)
 - The main part of the word is the so-called **stem**.
 - Other parts are called **affixes**.
 - E.g. plural dogs = morpheme “dog” (stem) + morpheme “s” (affix)
 - E.g. plural affixes = “affix” (stem) + “es” (affix)
- Normalization usually means reducing each word to its stem.
 - Therefore, lemmatization often is referred to as **stemming**.

Lemmatization Algorithms

- A problem with stemming, it is not particularly well defined.
- Many irregular cases
 - women → woman
 - lying → lie
- A rather old, but still used, algorithm for lemmatization is the **Porter stemmer**.
- Alternative: **Lancaster stemmer** (other rule set)
- Alternative: **WordNet lemmatizer**
 - Removes affixes only if the resulting word is in its vocabulary.
 - Slower performance due to this comparison.

Porter Stemmer

- Porter stemmer uses stemming rules it applies to words.
 - The rules are applied in sequence (5 steps).
- Examples for stemming rules:
 - Step 1: Plural and Participals
 - SSES -> SS | IES -> I | SS -> SS | S ->
 - caresses -> caress | ponies -> poni | caress -> caress | cats -> cat
 - Step 2
 - ATIONAL -> ATE relational -> relate | TIONAL -> TION conditional -> condition, rational -> ration
- Find more here: <https://tartarus.org/martin/PorterStemmer>

<https://tartarus.org/martin/PorterStemmer/def.txt>

Stemming – General

- Choose a stemmer that fits your application.
 - Do you search in the text and want to find alternative forms of search words?
 - In this case Porter stemmer is a valid choice.
 - Do you require valid words as lemmas?
 - In this case the WordNet lemmatizer is preferable.

Practical Part

- Please follow the section about stemming in Notebook NLP02



DISTANCES BETWEEN STRINGS

String Comparison

- **Comparing two strings** is a usual operation.
- Consider **“inention”**, which is obviously not a complete word but contains some error.
 - Maybe, the author wanted to write “intention”
 - A mere string comparison will fail “inention” == “intention” → False
 - Maybe the author meant to write “intervention”? Or “invention”? Or “mention”?
- What is the difference between the strings “inention”, “intention”, “intervention”, “invention”, “mention”, ...
- A more complicated example: “OTH president Prof. Bulitta” vs. “OTH president Bulitta”
 - Both strings mean the same person, however, they are not equal.

Edit Distance

- Edit Distance is a distance measure based on the number of elemental string operations to transform one string into another one.
- Elemental string operations:
 - Insertion
 - Deletion
 - Substitution
- Alignment: Correspondence between two substrings.
 - Needs to be considered for this kind of string comparison.

Alignment

Example

- Different alignments for our example “inention” and “invention”:

I	N	V	E	N	T	I	O	N
I	N	E	N	T	I	O	N	*
*	I	N	E	N	T	I	O	N
I	N	*	E	N	T	I	O	N

Edit Distance

- Levenshtein distance: Count the number of insertion and deletion steps to transform one string into another.
- No substitution operation, however, substitution = 1 deletion + 1 insertion
- Looking at the alignment: We can compute various distances between “inention” and “invention”

Alignment

Levenshtein distance

- Different alignments lead to different Levenshtein distances for our example “inention” and “invention”:

I	N	V	E	N	T	I	O	N	
I	N	E	N	T	I	O	N	*	13
*	I	N	E	N	T	I	O	N	5
I	N	*	E	N	T	I	O	N	1

- To determine the distance between two strings, we are looking for their **minimum edit distance**.
- Compare two strings of (not necessarily equal) lengths n and m .
 - Brute force search will not work for longer strings due to complexity
 - $n > m$: $\binom{n}{m}$ can get become a very large number, much larger than n .
- Efficient algorithm works with recurrence relation
 - Let $D(i,j)$ be the distance between $\text{string1}[0:i]$ and $\text{string2}[0:j]$

$$- D(i,j) = \min \left\{ \begin{array}{l} D(i-1,j) + \text{delete_cost}(\text{string1}[i]) \\ D(i,j-1) + \text{insert_cost}(\text{string2}[j]) \\ D(i-1,j-1) + \text{substitution_cost}(\text{string1}[i], \text{string2}[j]) \end{array} \right\}$$

Minimum Edit Distance Algorithm

```
function MIN-EDIT-DISTANCE(source, target) returns min-distance

   $n \leftarrow \text{LENGTH}(\textit{source})$ 
   $m \leftarrow \text{LENGTH}(\textit{target})$ 
  Create a distance matrix  $D[n+1, m+1]$ 

  # Initialization: the zeroth row and column is the distance from the empty string
   $D[0,0] = 0$ 
  for each row  $i$  from 1 to  $n$  do
     $D[i,0] \leftarrow D[i-1,0] + \textit{del-cost}(\textit{source}[i])$ 
  for each column  $j$  from 1 to  $m$  do
     $D[0,j] \leftarrow D[0,j-1] + \textit{ins-cost}(\textit{target}[j])$ 

  # Recurrence relation:
  for each row  $i$  from 1 to  $n$  do
    for each column  $j$  from 1 to  $m$  do
       $D[i,j] \leftarrow \text{MIN}( D[i-1,j] + \textit{del-cost}(\textit{source}[i]),$ 
                           $D[i-1,j-1] + \textit{sub-cost}(\textit{source}[i], \textit{target}[j]),$ 
                           $D[i,j-1] + \textit{ins-cost}(\textit{target}[j]))$ 

  # Termination
  return  $D[n,m]$ 
```

From the book of Jurafsky and Martin (Figure 2.17)

Literature

Zweite Headline

- Daniel Jurafsky, James H. Martin , Speech and Language Processing, Copyright © 2023. All rights reserved. Draft of January 7, 2023.
Online: https://web.stanford.edu/~jurafsky/slp3/ed3book_jan72023.pdf
- Steven Bird, Ewan Klein, and Edward Loper, Natural Language Processing with Python – Analyzing Text with the Natural Language Toolkit, O-Reilly, 3rd edition (2019), Online Version <https://www.nltk.org/book/>
- <https://docs.python.org/3/howto/regex.html#regex-howto>

Quellen: